



Chainlink LlamaRisk LlamaGuard NAV CRE

Security Assessment

February 27, 2026

Prepared for:

LlamaRisk

LlamaRisk

Prepared by: **Tarun Bansal**

Table of Contents

Table of Contents	1
Project Summary	2
Executive Summary	3
Project Goals	5
Project Targets	6
Project Coverage	7
Summary of Findings	8
Detailed Findings	9
1. Missing market authorization check allows unauthorized markets to use oracle data	9
2. Lack of workflow existence check when activating workflows	11
3. Missing event emission for oracle address update	12
A. Vulnerability Categories	13
B. Fix Review Results	15
Detailed Fix Review Results	15
C. Fix Review Status Categories	16
About Trail of Bits	17
Notices and Remarks	18

Project Summary

Contact Information

The following project manager was associated with this project:

Mary O'Brien, Project Manager
mary.obrien@trailofbits.com

The following engineering director was associated with this project:

Benjamin Samuels, Engineering Director, Blockchain
benjamin.samuels@trailofbits.com

The following consultants were associated with this project:

Tarun Bansal, Consultant
tarun.bansal@trailofbits.com

Project Timeline

The significant events and milestones of the project are listed below.

Date	Event
January 14, 2026	Pre-project kickoff call
January 23, 2026	Delivery of report draft
January 27, 2026	Report readout meeting
February 27, 2026	Delivery of report with fix review results

Executive Summary

Engagement Overview

LlamaRisk engaged Trail of Bits to review the security of the LlamaGuard NAV CRE actions smart contracts for Horizon lending pools.

The LlamaGuard system provides validated Net Asset Value (NAV) price feeds for DeFi assets used in Horizon lending protocol markets, with automated protective mechanisms that freeze markets when price bounds are breached. The system comprises several interconnected components: the `LlamaGuardOracle` contract implements the Chainlink `AggregatorV3` interface to store validated NAV prices and risk parameters; the `LlamaGuardOracleProxy` contract receives and validates reports from the CRE workflow; the `ParameterRegistry` stores risk parameters governing bound calculations; and the `HorizonAgentHub` coordinates automated freeze actions through `HorizonFreezeAgent` when the oracle signals a bound breach.

A team of one consultant conducted the review from January 19 to January 23, 2026, for a total of one engineer-week of effort. Our testing efforts focused on identifying security vulnerabilities that could affect the liveness and correctness of the system. With full access to source code, documentation, and deployment configurations, we performed static and dynamic testing of the codebase using automated and manual processes. The CRE workflow code and the Horizon lending pool smart contracts were out of scope for this engagement.

Observations and Impact

The review identified three findings across access controls, data validation, and auditing categories. The codebase demonstrates several positive security practices, including comprehensive role-based access control mechanisms and well-documented contract interfaces that facilitate security analysis.

The most notable finding (**TOB-CLGN-1**) concerns the `LlamaGuardOracle`'s market authorization system. While the contract implements infrastructure to restrict which markets can use oracle data through an authorized markets list, the `getLatestUpdateByParameterAndMarket` function does not validate market addresses against this list before returning risk parameter data. In a misconfigured deployment where the `HorizonAgentHub` registers an agent with incorrect market addresses, this could enable freeze actions on unintended markets if those markets mistakenly grant the necessary permissions to the freeze agent. The risk is constrained by the requirement that multiple configuration errors must align for exploitation to occur.

Two informational findings address operational visibility gaps. The `setWorkflowActive` function can activate non-existent workflow configurations without verification, potentially creating a confusing contract state. Additionally, the `setLlamaGuardOracle` function

modifies critical oracle routing without emitting events, thereby reducing monitoring systems' ability to detect configuration changes.

Recommendations

Based on the findings identified during the security review, Trail of Bits recommends that LlamaRisk team take the following steps:

- **Remediate the findings disclosed in this report.** While the identified issues require specific configuration errors to be exploited, they represent gaps in defense-in-depth protections that should be addressed. These findings should be resolved through direct code changes or as part of broader system hardening efforts.
- **Analyse and document risks of parameter updates triggering an unexpected market freeze.** The ParameterRegistry contract allows the updater to update the parameters used for calculating the asset price bounds. A sudden change in the risk parameters can result in a regular price movement being considered an out-of-bound price, and it can trigger a market freeze for that asset. Analyse the effect of parameter updates on the extrapolated asset price movement with a simulation to ensure that the parameter updates do not trigger an unexpected market freeze. Document these risks and guidelines in the updater operations documentation.

Finding Severities and Categories

The following tables provide the number of findings by severity and category.

EXPOSURE ANALYSIS

<i>Severity</i>	<i>Count</i>
High	0
Medium	0
Low	1
Informational	2
Undetermined	0

CATEGORY BREAKDOWN

<i>Category</i>	<i>Count</i>
Access Controls	1
Auditing and Logging	1
Data Validation	1

Project Goals

The engagement was scoped to provide a security assessment of the LlamaGuard NAV CRE actions. Specifically, we sought to answer the following non-exhaustive list of questions:

- Are CRE updates parsed correctly in the smart contracts?
- Can an attacker bypass the access control system?
- Are there any denial-of-service (DoS) vulnerabilities affecting the oracle and agent smart contracts?
- Can parameter updates break the system?
- Can the oracle data be used for/by unauthorized markets?
- Do all the state-changing functions emit sufficient events to facilitate system monitoring?
- Does any update lead to an inconsistent system state?

Project Targets

The engagement involved reviewing and testing the following target.

LlamaGuard oracle and agent smart contracts

Repository	https://github.com/llama-risk/llamaguard-contracts
Version	40b9ec824a818005e600d45f5e3229c57ad16dac
Type	Solidity
Platform	EVM

Project Coverage

This section provides an overview of the analysis coverage of the review, as determined by our high-level engagement goals. Our approaches included the following:

- **General.** We reviewed the entire codebase for common Solidity flaws, such as missing contract existence checks on low-level calls, issues with access controls, unimplemented functions, and memory manipulation issues in assembly code blocks. We used Slither to statically analyze the code.
- **LlamaGuardOracle contracts.** We reviewed the LlamaGuardOracleProxy and LlamaGuardOracle contracts to check the implementation of the asset price and risk parameter updates. We checked the enforcement of access control checks in the CRE workflow receiver function and reviewed the update input parsing and storage implementation for correctness. We reviewed the market authorization checks and role-based access control checks for writing updates. Finally, we checked if all the critical state-changing functions emit sufficient events for system monitoring.
- **HorizonAgentHub and agent contracts.** We reviewed the HorizonAgentHub and HorizonFreezeAgent contracts for the correct implementation of the action data and market state validation. We assessed whether a frozen market can get a freeze command again, whether unfreeze action is triggered automatically, and whether malicious users can block a freeze action. We also reviewed the correctness and effectiveness of the access control checks.

Coverage Limitations

Because of the time-boxed nature of testing work, it is common to encounter coverage limitations. During this project, we were unable to perform comprehensive testing of the following system elements, which may warrant further review:

- CRE workflow implementation
- Off-chain components

Summary of Findings

The table below summarizes the findings of the review, including details on type and severity.

ID	Title	Type	Severity
1	Missing market authorization check allows unauthorized markets to use oracle data	Access Controls	Low
2	Lack of workflow existence check when activating workflows	Data Validation	Informational
3	Missing event emission for oracle address update	Auditing and Logging	Informational

Detailed Findings

1. Missing market authorization check allows unauthorized markets to use oracle data

Severity: Low

Difficulty: High

Type: Access Controls

Finding ID: TOB-CLGN-1

Target: src/LlamaGuardOracle.sol

Description

The LlamaGuardOracle contract implements a market authorization system through the `_authorizedMarkets` mapping and provides `addAuthorizedMarket` and `removeAuthorizedMarket` functions to manage which markets can legitimately use the oracle's risk parameter data. However, the `getLatestUpdateByParameterAndMarket` function does not validate whether the provided market address is authorized before returning risk parameter data. Instead, it unconditionally overwrites the market field of the returned `RiskParameterUpdate` struct with the caller-supplied market address.

```
function _getLatestUpdateByParameterAndMarket(
    bytes32 updateTypeHash,
    address market
)
    internal
    view
    returns (RiskParameterUpdate memory)
{
    uint256 updateId = _latestUpdateIdByType[updateTypeHash];

    require(updateId != 0, InvalidUpdateId(updateId));

    RiskParameterUpdate memory update = updateHistory[updateId];
    // Rewrite the input market to the returned RiskParameterUpdate
    update.market = market; // No authorization check
    return update;
}
```

Figure 1.1: The function rewrites the market field without validating authorization.
(src/LlamaGuardOracle.sol#L322-L339)

This allows any contract or external caller querying the oracle to obtain risk parameter updates for arbitrary market addresses, including markets not on the authorized list. While the oracle stores global updates with `market = address(0)`, the returned data can be

associated with any market address through the rewriting mechanism. In a misconfigured system where the HorizonAgentHub registers an agent with an incorrect market address in its `allowedMarkets` array, the agent could trigger freeze actions on unauthorized markets if those markets' `PoolConfigurator` contracts grant the freeze agent the necessary `RISK_ADMIN` role.

Exploit Scenario

An operator misconfigures the HorizonAgentHub and registers an agent with an `allowedMarkets` array that includes a market address not intended to be protected by this oracle instance. When the oracle receives a freeze state update from the CRE workflow, the agent queries `getLatestUpdateByParameterAndMarket` with the misconfigured market address. The oracle returns the global risk parameter data with the market field rewritten to the unauthorized address. The agent validates the freeze state and, if the unauthorized market's `PoolConfigurator` has granted the `RISK_ADMIN` role to the agent, executes `setReserveFreeze` on that market. This freezes an unintended market based on risk parameters meant for a different market, potentially disrupting protocol operations and user access.

Recommendations

Short term, implement an authorization check in `_getLatestUpdateByParameterAndMarket` that validates the provided market address against `_authorizedMarkets` before returning the update.

Long term, improve the test suite to implement tests verifying the market authorization checks in the oracle smart contract functions.

2. Lack of workflow existence check when activating workflows

Severity: **Informational**

Difficulty: **High**

Type: Data Validation

Finding ID: TOB-CLGN-2

Target: `src/LlamaGuardOracleProxy.sol`

Description

The `setWorkflowActive` function allows the contract owner to enable or disable workflow configurations by setting the `isActive` flag for a given workflow ID. However, the function does not verify whether a workflow configuration exists for the provided workflow ID before modifying its active status. This means that an owner can call `setWorkflowActive` with an arbitrary, non-existent workflow ID, which creates an empty workflow configuration with only the `isActive` field set.

```
function setWorkflowActive(bytes32 workflowId, bool isActive) external onlyOwner {
    WorkflowConfig storage config = workflowConfigs[workflowId];
    config.isActive = isActive;

    emit WorkflowConfigUpdated(
        workflowId, config.expectedForwarder, config.expectedAuthor,
        config.expectedWorkflowName, isActive
    );
}
```

*Figure 2.1: The function does not validate workflow existence before activation.
([src/LlamaGuardOracleProxy.sol#L115-L122](#))*

While this does not create an immediate security vulnerability because the `onReport` function validates all workflow parameters, including the forwarder address, author, and workflow name, it can lead to a confusing contract state. The emitted `WorkflowConfigUpdated` event will contain zero values for `expectedForwarder`, `expectedAuthor`, and `expectedWorkflowName`, potentially misleading off-chain monitoring systems or creating false positives in security alerts.

Recommendations

Short term, add a validation check in `setWorkflowActive` that verifies at least one non-zero field exists in the workflow configuration before allowing activation.

3. Missing event emission for oracle address update

Severity: **Informational**

Difficulty: **Low**

Type: Auditing and Logging

Finding ID: TOB-CLGN-3

Target: `src/LlamaGuardOracleProxy.sol`

Description

The `setLlamaGuardOracle` function allows the contract owner to update the address of the underlying `LlamaGuardOracle` contract that receives validated reports. This is a critical configuration change that affects the destination of all incoming CRE workflow reports. However, the function does not emit an event after successfully updating the oracle address, making it difficult for off-chain systems to track when the oracle configuration changes.

```
function setLlamaGuardOracle(address newLlamaGuardOracle) external onlyOwner {
    ILlamaGuardOracle newLlamaguardOracle = ILlamaGuardOracle(newLlamaGuardOracle);
    require(newLlamaguardOracle.hasWriteAccess(address(this)),
InvalidLlamaGuardOracle());
    llamaguardOracle = newLlamaguardOracle;
    // No event emitted
}
```

*Figure 3.1: The function updates the oracle address without emitting an event
(`src/LlamaGuardOracleProxy.sol#L80-L84`)*

Without event emission, monitoring systems cannot detect oracle address changes through standard event logs. This reduces operational visibility and may delay detection of unauthorized or incorrect oracle updates. While the function includes proper access control through the `onlyOwner` modifier and validates that the new oracle grants write access to the proxy, the lack of event logging makes it harder to audit the contract's operational history and maintain comprehensive records of configuration changes.

Recommendations

Short term, add an event emission to `setLlamaGuardOracle` that logs both the previous and new oracle addresses whenever the configuration changes.

Long term, implement comprehensive event logging for all administrative functions that modify critical contract configurations, and establish a monitoring infrastructure that alerts on all configuration change events to enable rapid detection and response to unauthorized modifications.

A. Vulnerability Categories

The following tables describe the vulnerability categories, severity levels, and difficulty levels used in this document.

Vulnerability Categories	
Category	Description
Access Controls	Insufficient authorization or assessment of rights
Auditing and Logging	Insufficient auditing of actions or logging of problems
Authentication	Improper identification of users
Configuration	Misconfigured servers, devices, or software components
Cryptography	A breach of system confidentiality or integrity
Data Exposure	Exposure of sensitive information
Data Validation	Improper reliance on the structure or values of data
Denial of Service	A system failure with an availability impact
Error Reporting	Insecure or insufficient reporting of error conditions
Patching	Use of an outdated software package or library
Session Management	Improper identification of authenticated users
Testing	Insufficient test methodology or test coverage
Timing	Race conditions or other order-of-operations flaws
Undefined Behavior	Undefined behavior triggered within the system

Severity Levels	
Severity	Description
Informational	The issue does not pose an immediate risk but is relevant to security best practices.
Undetermined	The extent of the risk was not determined during this engagement.
Low	The risk is small or is not one the client has indicated is important.
Medium	User information is at risk; exploitation could pose reputational, legal, or moderate financial risks.
High	The flaw could affect numerous users and have serious reputational, legal, or financial implications.

Difficulty Levels	
Difficulty	Description
Undetermined	The difficulty of exploitation was not determined during this engagement.
Low	The flaw is well known; public tools for its exploitation exist or can be scripted.
Medium	An attacker must write an exploit or will need in-depth knowledge of the system.
High	An attacker must have privileged access to the system, may need to know complex technical details, or must discover other weaknesses to exploit this issue.

B. Fix Review Results

When undertaking a fix review, Trail of Bits reviews the fixes implemented for issues identified in the original report. This work involves a review of specific areas of the source code and system configuration, not a comprehensive analysis of the system.

On February 2, 2026, Trail of Bits reviewed the fixes and mitigations implemented by the LlamaRisk team for the issues identified in this report. We reviewed each fix to determine its effectiveness in resolving the associated issue.

In summary, of the three issues described in this report, LlamaRisk has resolved three issues. For additional information, please see the Detailed Fix Review Results below.

ID	Title	Severity	Status
1	Missing market authorization check allows unauthorized markets to use oracle data	Low	Resolved
2	Lack of workflow existence check when activating workflows	Informational	Resolved
3	Missing event emission for oracle address update	Informational	Resolved

Detailed Fix Review Results

TOB-CLGN-1: Missing market authorization check allows unauthorized markets to use oracle data

Resolved in [PR #19](#). The `_getLatestUpdateByParameterAndMarket` function now validates that the provided market is an authorized market.

TOB-CLGN-2: Lack of workflow existence check when activating workflows

Resolved in [PR #19](#). The `setWorkflowActive` function now validates that the provided `workflowId` corresponds to an existing workflow by checking the workflow config values.

TOB-CLGN-3: Missing event emission for oracle address update

Resolved in [PR #19](#). The `setLlamaGuardOracle` function now emits the `LlamaGuardOracleUpdated` event.

C. Fix Review Status Categories

The following table describes the statuses used to indicate whether an issue has been sufficiently addressed.

Fix Status	
Status	Description
Undetermined	The status of the issue was not determined during this engagement.
Unresolved	The issue persists and has not been resolved.
Partially Resolved	The issue persists but has been partially resolved.
Resolved	The issue has been sufficiently resolved.

About Trail of Bits

Founded in 2012 and headquartered in New York, Trail of Bits provides technical security assessment and advisory services to some of the world's most targeted organizations. We combine high-end security research with a real-world attacker mentality to reduce risk and fortify code. With 100+ employees around the globe, we've helped secure critical software elements that support billions of end users, including Kubernetes and the Linux kernel.

We maintain an exhaustive list of publications at <https://github.com/trailofbits/publications>, with links to papers, presentations, public audit reports, and podcast appearances.

In recent years, Trail of Bits consultants have showcased cutting-edge research through presentations at CanSecWest, HCSS, Devcon, Empire Hacking, GrrCon, LangSec, NorthSec, the O'Reilly Security Conference, PyCon, REcon, Security BSides, and SummerCon.

We specialize in software testing and code review assessments, supporting client organizations in the technology, defense, blockchain, and finance industries, as well as government entities. Notable clients include HashiCorp, Google, Microsoft, Western Digital, Uniswap, Solana, Ethereum Foundation, Linux Foundation, and Zoom.

To keep up with our latest news and announcements, please follow [@trailofbits on X](#) or [LinkedIn](#) and explore our public repositories at <https://github.com/trailofbits>. To engage us directly, visit our "Contact" page at <https://www.trailofbits.com/contact> or email us at info@trailofbits.com.

Trail of Bits, Inc.

228 Park Ave S #80688

New York, NY 10003

<https://www.trailofbits.com>

info@trailofbits.com

Notices and Remarks

Copyright and Distribution

© 2026 by Trail of Bits, Inc.

All rights reserved. Trail of Bits hereby asserts its right to be identified as the creator of this report in the United Kingdom.

Trail of Bits considers this report public information; it is licensed to LlamaRisk under the terms of the project statement of work and has been made public at LlamaRisk's request. Material within this report may not be reproduced or distributed in part or in whole without Trail of Bits' express written permission.

The sole canonical source for Trail of Bits publications is the [Trail of Bits Publications page](#). Reports accessed through sources other than that page may have been modified and should not be considered authentic.

Test Coverage Disclaimer

Trail of Bits performed all activities associated with this project in accordance with a statement of work and an agreed-upon project plan.

Security assessment projects are time-boxed and often rely on information provided by a client, its affiliates, or its partners. As a result, the findings documented in this report should not be considered a comprehensive list of security issues, flaws, or defects in the target system or codebase.

Trail of Bits uses automated testing techniques to rapidly test software controls and security properties. These techniques augment our manual security review work, but each has its limitations. For example, a tool may not generate a random edge case that violates a property or may not fully complete its analysis during the allotted time. A project's time and resource constraints also limit their use.